Actual training and evaluation: complete CPU episodes, CUDA PPO, frozen Actor testing

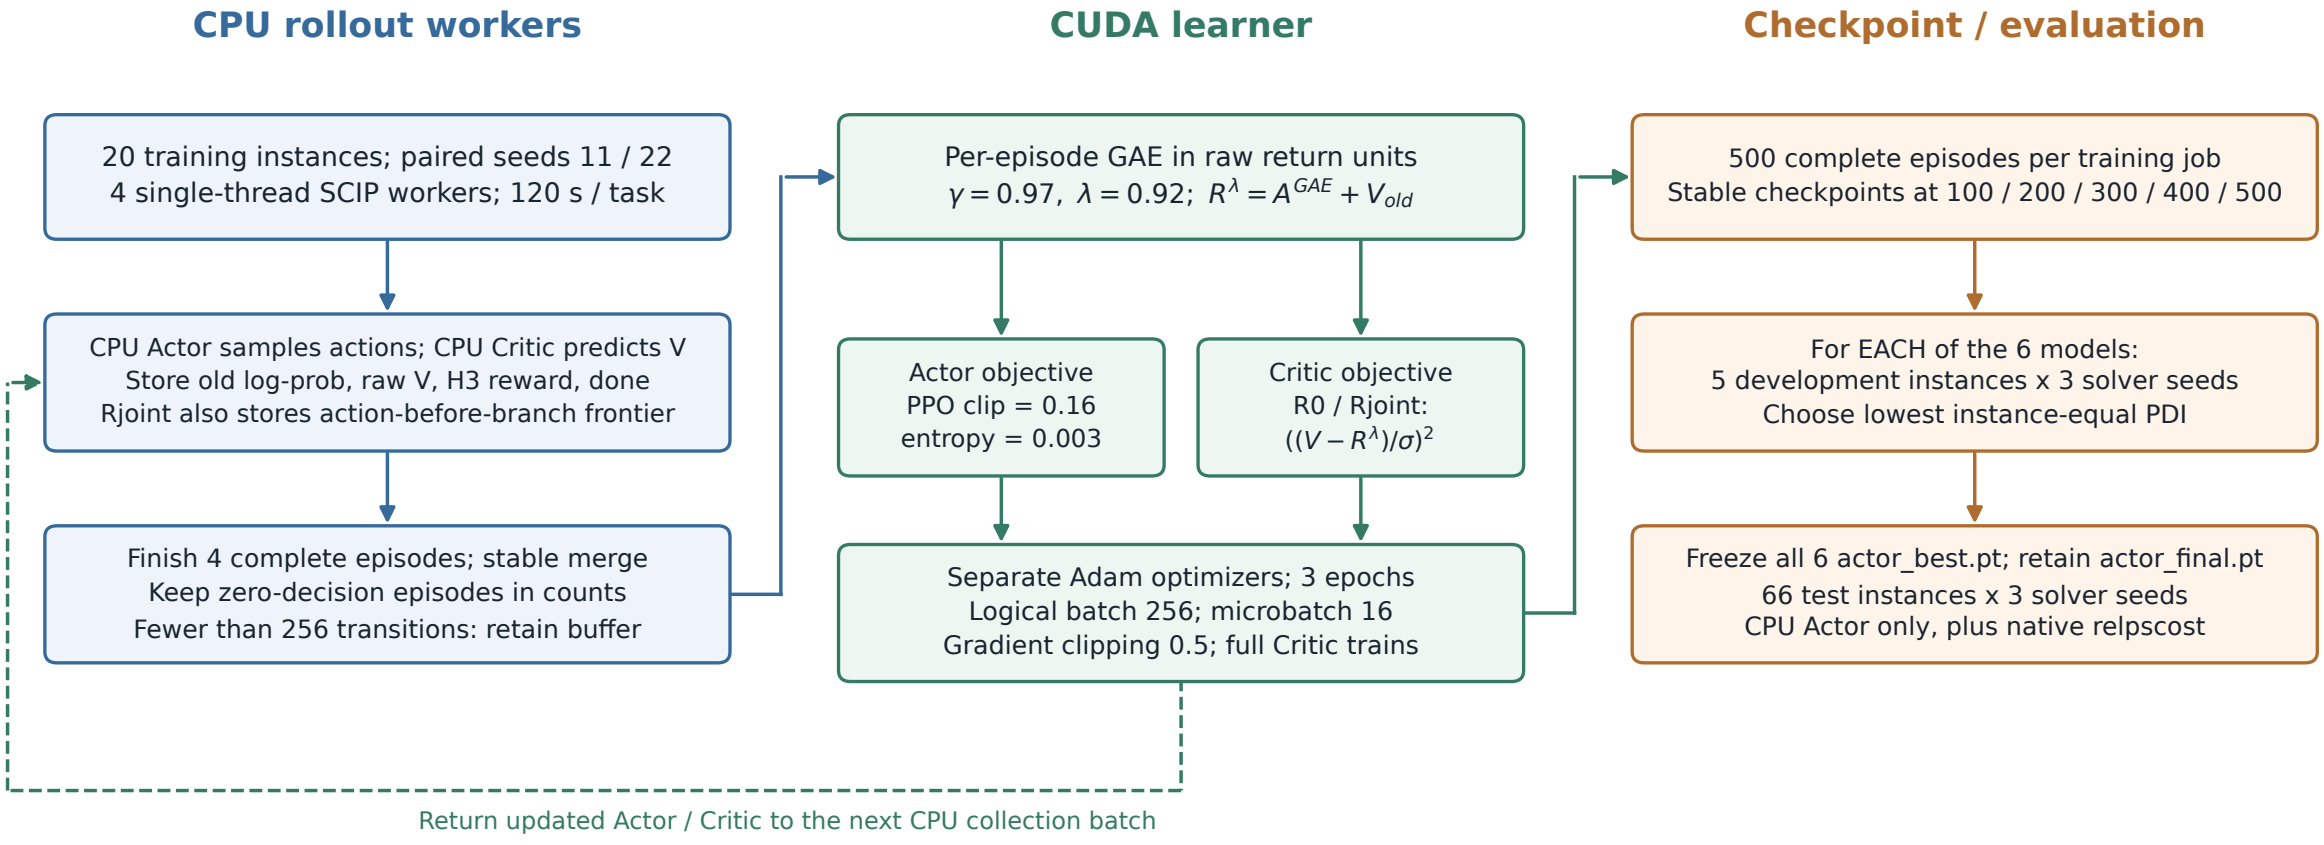

Offline diagnostic is separate: fixed Actor, shared Monte Carlo return labels, frozen encoder, value-head / context fitting.

Formal online models start fresh; only train-only calibration statistics are reused. No offline-head warm start and no test-based selection.

Evaluation disables Critic, frontier extraction, reward computation and PPO updates.