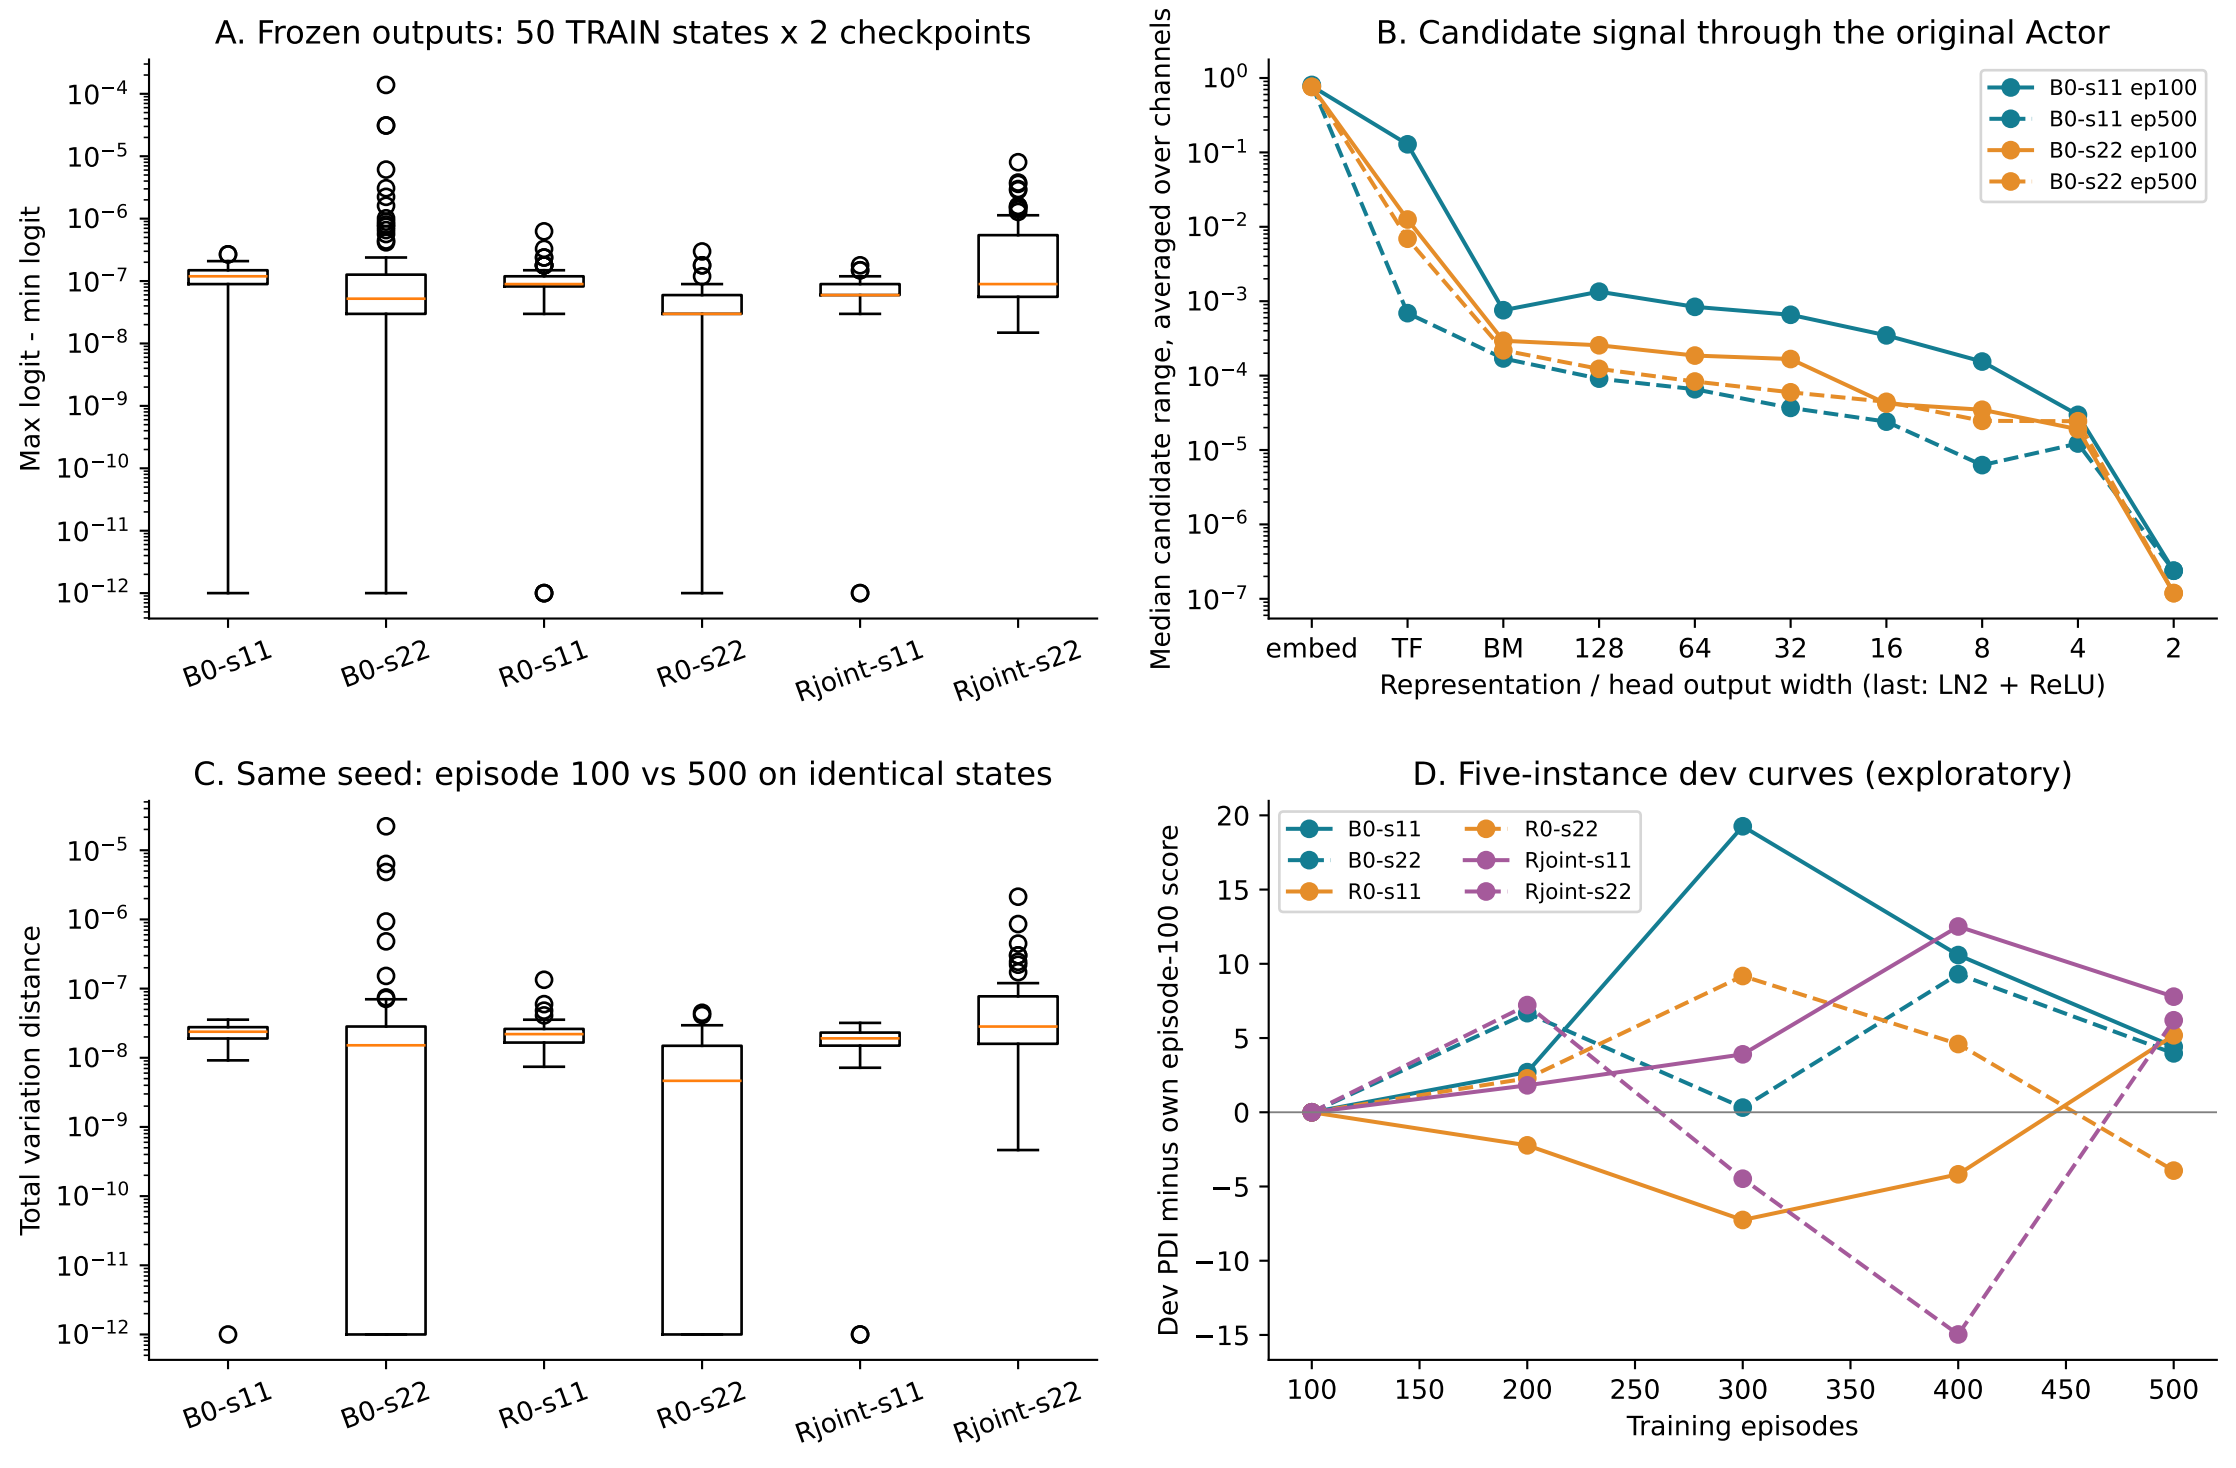
A. Frozen outputs: 50 TRAIN states x 2 checkpoints
Max logit - min logit
B0-s11
B0-s22
R0-s11
R0-s22
Rjoint-s11
Rjoint-s22
B. Candidate signal through the original Actor
Median candidate range, averaged over channels
B0-s11 ep100
B0-s11 ep500
B0-s22 ep100
B0-s22 ep500
embed
TF
BM
128
64
32
16
8
4
2
Representation / head output width (last: LN2 + ReLU)
C. Same seed: episode 100 vs 500 on identical states
Total variation distance
D. Five-instance dev curves (exploratory)
Dev PDI minus own episode-100 score
Training episodes